

R Workshop: Correlation

Brier Gallihugh, M.S.

2024-02-23

Table of contents

Introduction	1
Statistical Assumptions	2
Normality of Variables	2
Graphical Depiction of Normality Assumption (mpg)	2
Statistical Depiction of Normality Assumption (mpg)	3
Graphical Depiction of Normality Assumption (wt)	3
Statistical Depiction of Normality Assumption (wt)	4
Running An Actual Correlation	4

Introduction

Correlations might be the most common statistical test run (especially early on in a project). For this part of the workshop we will be using the `mtcars` data set due to its many numerical values that we can assess for correlation. Below we will see the first 10 rows of the data set displayed

```
library(tidyverse)
library(car)
library(psych)

data <- mtcars

print(head(data,10))
```

①

- ① Display the first 10 rows of the `mtcars` data. You can also return the last 10 rows with the following function `tail(data,10)`

	mpg	cyl	displacement	horsepower	drat	weight	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

Statistical Assumptions

The primary assumption of a **Pearson's correlation coefficient** is that the data is on some kind of interval scale. However, if we wish to generalize, we must have a random large sample (unlikely) and the individual variables should be roughly normally distributed. This is the assumption we will focus on for this part of the workshop.

Normality of Variables

For this part of the workshop, we are going to focus on the `mpg` and `wt` variables. We will be looking at normality both statistically₁ as well as graphically.

! A Note About Statistical Assumption Testing


1. Odds are your statistical test is going to fail pretty much every single time. Particularly if you use something like a Shapiro Wilk test, but we'll look at it anyway

Graphical Depiction of Normality Assumption (mpg)

```
mpg_plot <- ggplot(data,aes(x = mpg)) +
  geom_histogram(aes(y=after_stat(density))) +
  stat_function(fun = dnorm,
               args = list(mean = mean(data$mpg),
                           sd = sd(data$mpg)),
               col = "blue") +
  theme_classic()

print(mpg_plot)
```

- ① This should look very familiar to the histogram part of the workshop



Statistical Depiction of Normality Assumption (mpg)

```
print(psych::describe(data$mpg))
```

①

```
print(shapiro.test(data$mpg))
```

②

- ① The `psych` package has a bunch of nifty functions for social science research. One is the `describe()` function which gives you a bunch of variable level summary statistics (e.g., mean, median, se, etc.)

- ② The `shapiro.test()` performs a Shapiro Wilk test of normality. Keep in mind this particular test is very sensitive to sample sizes

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X1	1	32	20.09	6.03	19.2	19.7	5.41	10.4	33.9	23.5	0.61	-0.37	1.07

Shapiro-Wilk normality test


data: data\$mpg

W = 0.94756, p-value = 0.1229

Graphical Depiction of Normality Assumption (wt)

```
wt_plot <- ggplot(data, aes(x = wt)) +
  geom_histogram(aes(y=after_stat(density))) +
  stat_function(fun = dnorm,
               args = list(mean = mean(data$wt),
                           sd = sd(data$wt)),
               col = "blue") +
  theme_classic()

print(wt_plot)
```



Statistical Depiction of Normality Assumption (wt)

```
print(psych::describe(data$wt))
print(shapiro.test(data$wt))
```

```
vars  n mean  sd median trimmed  mad  min  max range skew kurtosis  se
X1    1 32 3.22 0.98  3.33   3.15 0.77 1.51 5.42  3.91 0.42   -0.02 0.17
Shapiro-Wilk normality test
```

```
data: data$wt
W = 0.94326, p-value = 0.09265
```

Running An Actual Correlation

There are multiple packages and methods for calculating a correlation in R depending on what you want to assess. The best to use for psychology is probably the `corr.test()` function in the `psych` package because it allows you to change the type of correlation you wish to compute (e.g., spearman vs pearson) as well as generate confidence intervals and do p value adjustments

```
corr_results <- corr.test(x = data$mpg, ①
                           y = data$wt, ②
```

```
use = "pairwise",
method = "pearson",
adjust = "holm")
```

③
④
⑤

```
?corr.test()
```

- ① Choose one of your variables to be your x variable
- ② Choose the other to be your y variable
- ③ You can choose “pairwise” or “complete”. For information on what each does, use the following function to access the documentation: `?psych::corr.test()`
- ④ You can adjust method to be other ones like “spearman”
- ⑤ You can also use “bonferroni” among a few others

Below we will see the output of the correlation results as you might be used to seeing in a program like SPSS.

```
print(corr_results)
```

```
Call:corr.test(x = data$mpg, y = data$wt, use = "pairwise", method = "pearson",
  adjust = "holm")
```

Correlation matrix

```
[1] -0.87
```

Sample Size

```
[1] 32
```

These are the unadjusted probability values.

The probability values adjusted for multiple tests are in the p.adj object.
[1] 0

To see confidence intervals of the correlations, print with the `short=FALSE` option

While the above is great, notice we didn’t get a confidence interval output despite asking for it with `ci = TRUE`. Sometimes R will store complex computations within the output object (e.g., `corr_results`). To get this output we can put a `$` after the output. If there is extra information stored, but not shown, we’ll get a drop down box. We want the `ci` option. Below we will see the output that results from this. We should see the following:

$r = -0.87$, $p < .001$ with a CI = [-.93,-.74]

```
print(corr_results$ci)
```

	lower	r	upper	p
NA-NA	-0.9338264	-0.8676594	-0.7440872	1.293959e-10